

# Model 1: SVM

Junaira I. Ibrahim

2022-12-16

## *Importing Packages*

```
# Helper packages
```

```
library(dplyr) # for data wrangling
```

```
## Warning: package 'dplyr' was built under R version 4.1.3
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(ggplot2) # for awesome graphics
```

```
## Warning: package 'ggplot2' was built under R version 4.1.3
```

```
library(rsample) # for data splitting
```

```
## Warning: package 'rsample' was built under R version 4.1.3
```

```
library(caret) # for classification and regression training
```

```
## Warning: package 'caret' was built under R version 4.1.3
```

```
## Loading required package: lattice
```

```
library(kernlab) # for fitting SVMs
```

```
## Warning: package 'kernlab' was built under R version 4.1.3
```

```
##
```

```
## Attaching package: 'kernlab'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## alpha
```

```
library(modeldata) #for Failure.binary data
```

```
## Warning: package 'modeldata' was built under R version 4.1.3
```

```
library(forcats)
```

```
## Warning: package 'forcats' was built under R version 4.1.3
```

```
library(bestNormalize)
```

```
## Warning: package 'bestNormalize' was built under R version 4.1.3
```

```
library(pdp)      # for partial dependence plots, etc.
```

```
## Warning: package 'pdp' was built under R version 4.1.3
```

```
library(vip)      # for variable importance plots
```

```
## Warning: package 'vip' was built under R version 4.1.3
```

```
##
```

```
## Attaching package: 'vip'
```

```
## The following object is masked from 'package:utils':
```

```
##
```

```
##      vi
```

**Importing the dataset** The dataset used in this model is imported from `radiomics` data. It has 197 observations and 431 variables.

```
datard <- read.csv("D:/MS_STATISTICS/STT225 Statistical Computing/FINAL PROJECT/radiomics_completedata.csv")
dim(datard)
```

```
## [1] 197 431
```

```
#####Checking for null and missing values
```

```
is.na(datard)
```

```
colSums(is.na(datard)) # no NA thus, there is no missing values
```

```
#####Checking for normality
```

```
md1 = datard%>%select_if(is.numeric)
```

```
datamd1 = lapply(md1[, -1], shapiro.test)
```

```
r = lapply(datamd1, function(x)x$p.value) #Extracting p-value only
```

```
s=unlist(r)      #to convert a list to vector
```

```
sum(s[s>0.05])
```

```
## [1] 0.1350113
```

```
r$Entropy_cooc.W.ADC
```

```
## [1] 0.1350113
```

Based on the results, there is only one variable who is normally distributed (i.e. `Entropy_cooc.W.ADC`). All the rest are not normally distributed. Hence, we will try to normalize the data using `orderNorm()` function.

```
#####Normalizing the data
```

```
datard_norm = datard[,c(3,5:length(names(datard)))]
```

```
datard_norm = apply(datard_norm,2,orderNorm)
```

```
datard_norm = lapply(datard_norm, function(x) x$x.t) #to transformed original data
```

```
datard_norm = datard_norm%>%as.data.frame()
```

Check the new data for normality

```
data1r2 = lapply(datard_norm, shapiro.test)
```

```
r2 = lapply(data1r2, function(x) x$p.value)
```

```
s2 = unlist(r2)
```

```
sum(s2>0.05)
```

```
## [1] 428
```

Based on the results, the rest of the variables is now normally distributed.

Substituting the normalized values into the original data, we have

```
r3 = select(datard, c("Failure.binary", "Entropy_cooc.W.ADC"))
datard_n = cbind(r3, datard_norm)
```


To set the Failure.binary into a factor level in dataset datard\_n, we use the function as.factor() function.

```
datard_n$Failure.binary=as.factor(datard_n$Failure.binary)
```

In this session, we will use datard\_n.

Getting the correlation of the whole data

```
#correlation
newdatard = select(datard_n, -c("Failure.binary"))
cor.newdatard = cor(newdatard)
corr = round(cor.newdatard, 2) # 2 decimals
heatmap(corr, Rowv=NA, Colv=NA, scale="none", revC = T)
```



#Support Vector Machine Support vector machines (SVMs) offer a direct approach to binary classification.

SVMs use the kernel trick to enlarge the feature space using basis functions. A **Kernel Trick** is a simple method where a Non Linear data is projected onto a higher dimension space so as to make it easier to classify the data where it could be linearly divided by a plane. The popular kernel function used by SVMs are Linear "svmLinear", Polynomial Kernel "svmPoly" and Radial basis kernel "svmRadial".

In the following chunks, we use getModelInfo() function to extract the hyperparameters from various SVM

implementations with different kernel functions.

```
# Linear (i.e., soft margin classifier)
caret::getModelInfo("svmLinear")$svmLinear$parameters
```

```
##   parameter  class label
## 1          C numeric  Cost
```

```
# Polynomial kernel
caret::getModelInfo("svmPoly")$svmPoly$parameters
```

```
##   parameter  class      label
## 1    degree numeric Polynomial Degree
## 2      scale numeric      Scale
## 3          C numeric      Cost
```

```
# Radial basis kernel
caret::getModelInfo("svmRadial")$svmRadial$parameters
```

```
##   parameter  class label
## 1      sigma numeric Sigma
## 2          C numeric  Cost
```

We can tune an SVM model with `train()` function with radial basis kernel using the data `rdtrain` and 10-fold CV.

The data `datard_n` is split into 80% of training data and 20% testing data.

```
#80% training data - 20% testing data
rdsplit <- initial_split(datard_n, prop = 0.8, strata = "Failure.binary")
rdsplit
```


```
## <Training/Testing/Total>
## <157/40/197>
```

```
rdtrain <- training(rdsplit)
rdtest  <- testing(rdsplit)
```

```
set.seed(1854) # for reproducibility
split_svm <- train(
  Failure.binary ~ .,
  data = rdtrain,
  method = "svmRadial",
  preProcess = c("center", "scale"),
  trControl = trainControl(method = "cv", number = 10),
  tuneLength = 10
)
```

Plotting the results, we see that smaller values of the cost parameter ( $C = 16-64$ ) provide better cross-validated accuracy scores for these training data.

```
# Plot results
ggplot(split_svm) + theme_light()
```



```
# Print results
split_svm$results
```

| ##    | sigma       | C      | Accuracy  | Kappa     | AccuracySD | KappaSD   |
|-------|-------------|--------|-----------|-----------|------------|-----------|
| ## 1  | 0.001838272 | 0.25   | 0.6627451 | 0.0000000 | 0.01891300 | 0.0000000 |
| ## 2  | 0.001838272 | 0.50   | 0.7023284 | 0.1640337 | 0.05605257 | 0.1606231 |
| ## 3  | 0.001838272 | 1.00   | 0.7407108 | 0.3543337 | 0.07178044 | 0.1596734 |
| ## 4  | 0.001838272 | 2.00   | 0.8018382 | 0.5274765 | 0.07671419 | 0.1924831 |
| ## 5  | 0.001838272 | 4.00   | 0.7955882 | 0.5235397 | 0.08028357 | 0.2025178 |
| ## 6  | 0.001838272 | 8.00   | 0.7826716 | 0.4985725 | 0.09350611 | 0.2356203 |
| ## 7  | 0.001838272 | 16.00  | 0.7901225 | 0.5162573 | 0.10016573 | 0.2493108 |
| ## 8  | 0.001838272 | 32.00  | 0.7963725 | 0.5359780 | 0.08950325 | 0.2136244 |
| ## 9  | 0.001838272 | 64.00  | 0.7901225 | 0.5191755 | 0.09042933 | 0.2158313 |
| ## 10 | 0.001838272 | 128.00 | 0.7963725 | 0.5359780 | 0.08950325 | 0.2136244 |

Control parameter

In order to obtain predicted class probabilities from an SVM, additional parameters need to be estimated. The predicted class probabilities are often more useful than the predicted class labels. For instance, we would need the predicted class probabilities if we were using an optimization metric like AUC. In that case, we can set `classProbs = TRUE` in the call to `trainControl()`.

```
class.weights = c("No" = 1, "Yes" = 10)
```

```
# Control params for SVM
ctrl <- trainControl(
  method = "cv",
  number = 10,
```

```

classProbs = TRUE,
summaryFunction = twoClassSummary
)

rdtrain$Failure.binary = fct_recode(rdtrain$Failure.binary,No="0",Yes="1")

```

Print the AUC values during Training

```

# Tune an SVM
set.seed(5628) # for reproducibility
train_svm_auc <- train(
  Failure.binary ~ .,
  data = rdtrain,
  method = "svmRadial",
  preProcess = c("center", "scale"),
  metric = "ROC",
  trControl = ctrl,
  tuneLength = 10
)

# Print results
train_svm_auc$results

```

```

##          sigma      C      ROC      Sens      Spec      ROCSD      SensSD
## 1  0.001628158  0.25 0.7606061 0.8736364 0.4566667 0.12607169 0.09422732
## 2  0.001628158  0.50 0.7606061 0.8336364 0.4766667 0.12607169 0.17726884
## 3  0.001628158  1.00 0.7934545 0.8936364 0.4566667 0.12532974 0.07184439
## 4  0.001628158  2.00 0.8294545 0.9036364 0.5366667 0.09189053 0.06388120
## 5  0.001628158  4.00 0.8288182 0.8827273 0.5533333 0.09509364 0.12192593
## 6  0.001628158  8.00 0.8236364 0.8836364 0.5533333 0.09767431 0.12790607
## 7  0.001628158 16.00 0.8369697 0.8845455 0.5333333 0.10611360 0.09866417
## 8  0.001628158 32.00 0.8446364 0.8845455 0.5700000 0.11019957 0.09866417
## 9  0.001628158 64.00 0.8413030 0.8745455 0.6266667 0.11637329 0.12180160
## 10 0.001628158 128.00 0.8563030 0.9045455 0.6866667 0.11635923 0.08806146
##          SpecSD
## 1  0.2553381
## 2  0.2211362
## 3  0.2553381
## 4  0.2191172
## 5  0.2470567
## 6  0.2470567
## 7  0.2509242
## 8  0.2301100
## 9  0.2153378
## 10 0.1596292

```

```

confusionMatrix(train_svm_auc)

```

```

## Cross-Validated (10 fold) Confusion Matrix
##
## (entries are percentual average cell counts across resamples)
##
##          Reference
## Prediction  No  Yes
##          No  59.9 10.8

```

```
##          Yes  6.4 22.9
##
## Accuracy (average) : 0.828
```

Based on the result, it is clear that we do a far better job at predicting the Nos.

Print the AUC values during Testing

```
rdtest$Failure.binary = fct_recode(rdtest$Failure.binary, No="0", Yes="1")
```

```
# Tune an SVM with radial
set.seed(5628) # for reproducibility
test_svm_auc <- train(
  Failure.binary ~ .,
  data = rdtest,
  method = "svmRadial",
  preProcess = c("center", "scale"),
  metric = "ROC", # area under ROC curve (AUC)
  trControl = ctrl,
  tuneLength = 10
)
```

```
# Print results
test_svm_auc$results
```

```
##          sigma      C      ROC      Sens Spec      ROCSD      SensSD      SpecSD
## 1  0.001420254  0.25 0.7833333 0.9333333 0.30 0.3147603 0.2108185 0.4216370
## 2  0.001420254  0.50 0.7833333 0.8000000 0.55 0.3147603 0.2698880 0.4972145
## 3  0.001420254  1.00 0.6833333 0.8166667 0.15 0.3884919 0.2539807 0.3374743
## 4  0.001420254  2.00 0.7833333 0.9000000 0.35 0.3147603 0.2249829 0.4116363
## 5  0.001420254  4.00 0.6583333 0.9000000 0.30 0.3736085 0.2249829 0.4216370
## 6  0.001420254  8.00 0.7833333 0.9000000 0.35 0.3147603 0.2249829 0.4743416
## 7  0.001420254 16.00 0.7833333 0.9333333 0.50 0.3147603 0.1405457 0.4714045
## 8  0.001420254 32.00 0.7833333 0.9000000 0.35 0.3147603 0.2249829 0.4743416
## 9  0.001420254 64.00 0.7833333 0.9000000 0.35 0.3147603 0.2249829 0.4743416
## 10 0.001420254 128.00 0.7833333 0.9000000 0.45 0.3147603 0.2249829 0.4972145
```

```
confusionMatrix(test_svm_auc)
```

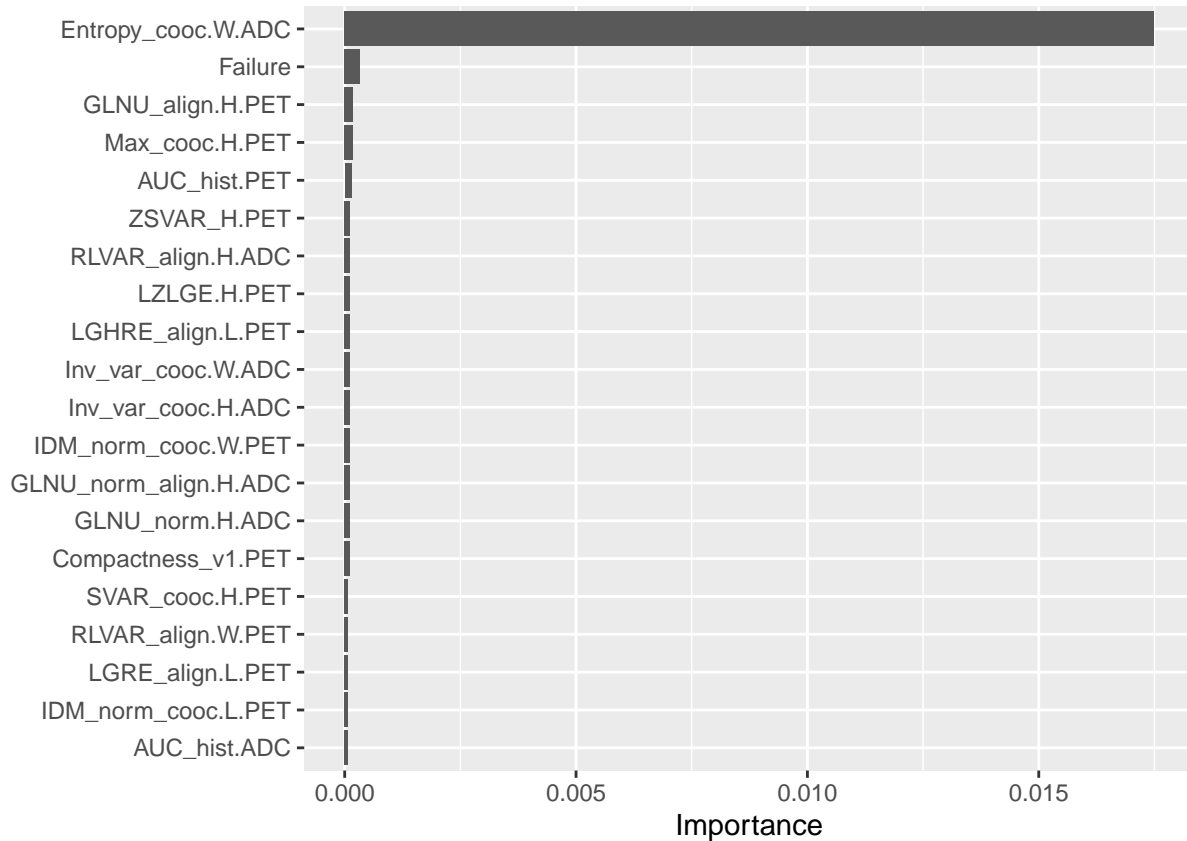
```
## Cross-Validated (10 fold) Confusion Matrix
##
## (entries are percentual average cell counts across resamples)
##
##          Reference
## Prediction  No  Yes
##          No 60.0 22.5
##          Yes 5.0 12.5
##
## Accuracy (average) : 0.725
```

Similar to training set, it is clear that we do a far better job at predicting the Nos.

To compute the vip scores we just call `vip()` with `method = "permute"` and pass our previously defined predictions wrapper to the `pred_wrapper` argument.

```
prob1 <- function(object, newdata) {
  predict(object, newdata = newdata, type = "prob")[, "Yes"]
}
```

```
# Variable importance plot
set.seed(2827) # for reproducibility
vip(train_svm_auc, method = "permute", nsim = 5, train = rdtrain, num_features = 20, target = "Failure.binary",
     pred_wrapper = prob1)
```



The results indicate that *Failure* and *Entropy\_cooc.W.ADC* is the most important feature in predicting *Failure.binary*.

Next, we use the *pdp* package to construct PDPs for the top four features according to the permutation-based variable importance scores.

```
features1 <- c("Failure", "Entropy_cooc.W.ADC",
               "RLVAR_align.H.ADC", "Compactness_v1.PET")
pdps <- lapply(features1, function(x) {
  partial(train_svm_auc, pred.var = x, which.class = 2,
          prob = TRUE, plot = TRUE, plot.engine = "ggplot2") +
  coord_flip()
})
grid.arrange(grobs = pdps, ncol = 2)
```



